



LCI Advanced Workshop 2025: Service Restoration & Repair

Alan Chapman
HPC Systems Analyst
Arizona State University Research Computing
alan.chapman@asu.edu

*This document is a result of work by volunteer LCI instructors and is licensed under CC BY-NC 4.0
(<https://creativecommons.org/licenses/by-nc/4.0/>)*

Agenda

- **Common Slurm Failures**
Resource allocation, job submission, MPI issues
- **Core Infrastructure Failures**
slurmctld, slurmdbd, slurmd daemon failures
- **Critical Recovery Procedures**
Emergency restoration, database recovery
- **Monitoring & Prevention**
Health checks, automated backups
- **Q&A & Discussion**

Diagnostic Workflow Overview

Standard Troubleshooting Sequence:

User-level diagnostics

<code>squeue -u \$USER -o "%i %t %r %S %M %l"</code>	<i># Job status</i>
<code>scontrol show job \$JOBID</code>	<i># Detailed job info</i>
<code>sacct -j \$JOBID --format=JobID,State,ExitCode</code>	<i># Accounting</i>
<code>tail -50 slurm-\$JOBID.out</code>	<i># Output logs</i>
<code>seff \$JOBID</code>	<i># Efficiency</i>

System-level diagnostics (head node)

<code>sinfo -N -l grep -v idle</code>	<i># Problem nodes</i>
<code>systemctl status slurmctld slurmdbd</code>	<i># Service status</i>
<code>journalctl -u slurmctld -n 50</code>	<i># System logs</i>

Tip: Always start with `scontrol show job $JOBID` – it reveals 80% of issues

Image source: <https://slurm.schedmd.com/network.html>

Resource Allocation Failures

Node Unavailable Issues:

Error: "Requested node configuration is not available"

Diagnosis

```
sinfo -N -l | grep -v idle  
scontrol show node lci-compute-40-1
```

Resolution

```
scontrol update NodeName=lci-compute-40-1 State=resume Reason="cleared"  
sbatch -x lci-compute-40-1,lci-compute-40-2 script.sh # Exclude problem nodes
```

QOS/Account Limits:

Error: "QOSMaxWallDurationPerJobLimit"

Check limits

```
sacctmgr show assoc user=$USER format=user,account,maxjobs,maxwall  
sprio -j $JOBID # Priority analysis
```

Fix

```
sbatch -t 23:59:00 -A different_account script.sh
```

Hardware Detection Failures

Memory Issues – OOM Kills

Error in slurmd.log: "Detected 1 oom-kill event(s)"

Diagnosis

```
sacct -j $JOBID --format=JobID,MaxRSS,ReqMem,Elapsed
```

```
dmesg | grep -i "killed process"
```

Resolution

```
sbatch --mem=64G script.sh # Increase memory
```

```
sbatch --mem-per-cpu=4G script.sh # Per-CPU allocation
```

GPU Problems

Error: "couldn't communicate with NVIDIA driver"

Diagnosis & Fix

```
srun --gres=gpu:1 nvidia-smi
```

```
sudo nvidia-smi -pm 1 # Enable persistence
```

```
sudo systemctl restart slurmd
```

```
scontrol update NodeName=node001 State=resume
```

Controller (slurmctld) Failures

Daemon Crash – Most Critical Failure

Symptoms

queue: Unable to contact slurm controller (connect failure)

Diagnosis

```
systemctl status slurmctld  
journalctl -u slurmctld -f  
tail -100 /var/log/slurm/slurmctld.log  
ss -tlnp | grep :6817
```

Recovery

```
systemctl start slurmctld
```

Debug mode if fails:

```
slurmctld -D -vvv
```

Controller (slurmctld) Failures

State File Corruption – Data Loss Risk

Error: "Unable to recover state from slurmctld.state"

Emergency Recovery (CAUSES JOB LOSS!)

```
systemctl stop slurmctld  
cd /var/lib/slurm  
cp slurmctld.state slurmctld.state.backup.$(date +%Y%m%d_%H%M)  
rm -f slurmctld.state  
systemctl start slurmctld
```

High Availability Failover

Controller Failover Procedure

```
# Primary controller failed  
# On backup controller:  
systemctl status slurmctld  
grep "PRIMARY\|BACKUP" /etc/slurm/slurm.conf  
  
# Force failover if backup doesn't auto-activate  
scontrol takeover  
  
# Repair and failback:  
# 1. Fix primary controller  
# 2. Update slurm.conf if needed  
# 3. scontrol reconfig on all nodes  
# 4. Let primary resume naturally
```

Best Practice: Test failover procedures regularly in maintenance windows

Database (slurmdbd) Failures

Database Daemon Issues

Symptoms: "SLURM accounting storage is disabled"

Diagnosis

```
systemctl status slurmdbd
```

```
tail -100 /var/log/slurm/slurmdbd.log
```

```
mysql -u slurm -p -h localhost -e "SHOW DATABASES;"
```

Recovery

```
systemctl start slurmdbd
```

Debug mode:

```
slurmdbd -D -vvv
```

Database Corruption

Error: "Table './slurm_acct_db/job_table' is marked as crashed"

Repair

```
systemctl stop slurmdbd
```

```
mysql -u slurm -p slurm_acct_db
```

```
CHECK TABLE job_table;
```

```
REPAIR TABLE job_table;
```

Mass repair:

```
mysqlcheck -u slurm -p --repair slurm_acct_db
```

Database Connection Issues

Connection Exhaustion

Error: "1040 Too many connections"

Diagnosis

```
mysql -u root -p -e "SHOW PROCESSLIST;"
```

```
mysql -u root -p -e "SHOW STATUS LIKE 'Threads_connected';"
```

Fix in /etc/mysql/mariadb.conf.d/50-server.cnf:

```
max_connections = 1000
```

```
wait_timeout = 28800
```

In slurmdbd.conf:

```
MaxQueryTimeRange=MONTH # Instead of INFINITE
```

```
systemctl restart mariadb slurmdbd
```

Compute Node (`slurmd`) Failures

Node Daemon Recovery

Node shows as "down" or "not responding"

Diagnosis

```
systemctl status slurmd
```

```
journalctl -u slurmd -n 50
```

```
tail -100 /var/log/slurm/slurmd.log
```

Recovery

```
systemctl start slurmd
```

```
scontrol update NodeName=node001 State=resume Reason="slurmd restarted"
```

Mass recovery:

```
pdsh -w node[001-100] "systemctl restart slurmd"
```

```
scontrol update NodeName=node[001-100] State=resume Reason="mass restart"
```

Hardware Configuration Mismatch

CPU/Memory Detection Issues

Error: "Node configuration differs from hardware: CPUs=64:128(hw)"

Diagnosis

```
scontrol show node node001 | grep -E "CPUs|RealMemory|Gres"
```

On compute node:

```
nproc
```

```
free -g
```

```
nvidia-smi -L
```

Resolution Options:

1. Update slurm.conf to match hardware

2. Use hardware detection

```
slurmd -C # Print actual hardware config
```

Best Practice: Use `slurmd -C` output to generate initial node configs

Authentication (Munge) Failures

Munge Key Issues - Security Critical

Error: "Munge credential decode failed: Invalid credential"

Diagnosis

`munge -n | unmunge` *# Test locally*

`pdsh -w node[001-100] "munge -n"` *# Test all nodes*

Recovery - Sync keys

`systemctl stop munge`

`pdcp -w node[001-100] /etc/munge/munge.key /etc/munge/`

`pdsh -w node[001-100] "chown munge:munge /etc/munge/munge.key"`

`pdsh -w node[001-100] "chmod 400 /etc/munge/munge.key"`

`pdsh -w node[001-100] "systemctl restart munge"`

Permission Problems

Fix munge directory permissions

`chown munge:munge /var/lib/munge /var/log/munge /run/munge`

`chmod 755 /var/lib/munge /var/log/munge /run/munge`

Network/Firewall Issues

Controller Unreachable

Diagnosis

```
ping slurmctld-host
telnet slurmctld-host 6817
ss -tlnp | grep 6817
iptables -L -n | grep 6817
```

Resolution - Open required ports

```
firewall-cmd --permanent --add-port=6817/tcp # slurmctld
firewall-cmd --permanent --add-port=6818/tcp # slurmd
firewall-cmd --permanent --add-port=6819/tcp # slurmdbd
firewall-cmd --reload
```

Critical Ports:

- 6817: slurmctld (controller)
- 6818: slurmd (compute nodes)
- 6819: slurmdbd (database)

Critical Recovery – Nuclear Option

Complete Cluster Recovery Procedure

When everything is broken - last resort

1. Stop all services

```
pdsh -w node[001-100] "systemctl stop slurmd"  
systemctl stop slurmctld slurmdbd
```

2. Backup critical state

```
cp -r /var/lib/slurm /backup/slurm_state_$(date +%Y%m%d_%H%M)  
mysqldump -u slurm -p slurm_acct_db > /backup/slurm_db_$(date +%Y%m%d).sql
```

3. Clean start (removes job state!)

```
rm -f /var/lib/slurm/slurmctld.state*
```

4. Start services in correct order

```
systemctl start slurmdbd  
systemctl start slurmctld  
pdsh -w node[001-100] "systemctl start slurmd"
```

5. Resume all nodes

```
scontrol update NodeName=ALL State=resume Reason="cluster restart"
```

Database Emergency Recovery

Recreate Corrupted Accounting Database

If accounting DB is completely corrupted

```
systemctl stop slurmdbd slurmctld
```

Recreate database

```
mysql -u root -p
```

```
DROP DATABASE slurm_acct_db;
```

```
CREATE DATABASE slurm_acct_db;
```

```
GRANT ALL ON slurm_acct_db.* TO 'slurm'@'localhost';
```

Reinitialize accounting

```
sacctmgr -i create cluster cluster_name
```

```
sacctmgr -i add account root Cluster=cluster_name
```

```
sacctmgr -i add user root DefaultAccount=root
```

```
systemctl start slurmdbd slurmctld
```


Configuration Synchronization

Config File Mismatch Issues

Error: "Node appears to have different slurm.conf"

Diagnosis

```
md5sum /etc/slurm/slurm.conf # Controller
```

```
pdsh -w node[001-100] "md5sum /etc/slurm/slurm.conf" # All nodes
```

Resolution

```
pdcp -w node[001-100] /etc/slurm/slurm.conf /etc/slurm/
```

```
pdsh -w node[001-100] "systemctl reload slurmd"
```

```
scontrol reconfig
```

Clock Synchronization

Error: "Message time stamp is too far in the future"

Check time sync

```
pdsh -w node[001-100] date
```

```
chrony sources -v
```

Force sync

```
pdsh -w node[001-100] "chrony makestep"
```

Monitoring & Health Checks

Automated Health Check Script

```
#!/bin/bash
# Controller health
if ! systemctl is-active slurmctld >/dev/null; then
    echo "CRITICAL: slurmctld is down"
    logger "SLURM: slurmctld service failed"
fi

# Database health
if ! systemctl is-active slurmdbd >/dev/null; then
    echo "CRITICAL: slurmdbd is down"
fi

# Node health
DOWN_NODES=$(sinfo -h -t down -o %N)
[ ! -z "$DOWN_NODES" ] && echo "WARNING: Down nodes: $DOWN_NODES"

# Database connectivity
if ! mysql -u slurm -p[password] -e "SELECT 1;" >/dev/null 2>&1; then
    echo "CRITICAL: Database connection failed"
fi
```

Deploy: Run every 5 minutes via cron, integrate with monitoring system

Backup Strategy

Automated Daily Backup

```
#!/bin/bash
DATE=$(date +%Y%m%d)
BACKUP_DIR="/backup/slurm"

# State files backup
mkdir -p $BACKUP_DIR/state
cp -r /var/lib/slurm/* $BACKUP_DIR/state/slurm_state_$DATE/

# Database backup
mysqldump -u slurm -p[password] slurm_acct_db | \
    gzip > $BACKUP_DIR/slurm_acct_db_$DATE.sql.gz

# Configuration backup
mkdir -p $BACKUP_DIR/config
cp /etc/slurm/* $BACKUP_DIR/config/slurm_config_$DATE/

# Cleanup (keep 30 days)
find $BACKUP_DIR -name "*" -mtime +30 -delete
```

Prevention Best Practices

Service Start Order

Always follow correct startup sequence:

1. `slurmdbd` # Database first
2. `slurmctld` # Controller second
3. `slurmd` # Compute nodes last

Regular Maintenance Tasks

- Weekly: Check for down nodes, review logs
- Monthly: Test backup/restore, failover procedures
- Quarterly: Update documentation, review capacity

Change Management

Before any config changes:

1. Backup current state
2. Test `in` development environment
3. Schedule maintenance window
4. Have rollback plan ready
5. Monitor post-change

Recovery Time Objectives

Recovery Time Objectives

Failure Type	RTO	Critical Steps
-----	-----	-----
Single node down	2-5 min	`systemctl restart slurmd`
Config mismatch	5-10 min	Sync configs, `scontrol reconfig`
slurmctld crash	1-3 min	`systemctl start slurmctld`
State corruption	15-30 min	Restore from backup
Database corruption	30-60 min	Repair tables or restore DB
Complete cluster failure	2-4 hours	Full recovery procedure

Escalation Path

1. L1: Basic service restarts, node resume
2. L2: Config sync, log analysis, backup restore
3. L3: Database recovery, network troubleshooting
4. Vendor: Hardware failures, software bugs

Common Pitfalls & Gotchas

What NOT to Do

- ✗ **Never** delete state files without backup
- ✗ **Never** restart services without checking dependencies
- ✗ **Never** modify database directly without stopping slurmdbd
- ✗ **Never** ignore authentication (munge) errors

Emergency Shortcuts That Backfire

DON'T do these under pressure:

```
rm -f /var/lib/slurm/*           # Loses all jobs
systemctl restart slurmctld slurmdbd  # Wrong order
scontrol update NodeName=ALL State=down  # Mass outage
```

Always Remember

- ✓ Backup before changes
- ✓ Check logs first
- ✓ Follow service dependencies
- ✓ Test in non-production first

Advanced Troubleshooting Tools

Log Analysis Commands

Real-time monitoring

```
journalctl -u slurmctld -f  
tail -f /var/log/slurm/slurmctld.log
```

Pattern matching

```
grep -E "(ERROR|FATAL|WARNING)" /var/log/slurm/*.log  
zgrep "node001" /var/log/slurm/slurmctld.log*
```

Performance analysis

```
strace -p $(pgrep slurmctld)  
perf top -p $(pgrep slurmctld)
```

Database Analysis

Connection monitoring

```
mysql -u root -p -e "SHOW PROCESSLIST;" | grep slurm
```

Query analysis

```
mysql -u root -p -e "SHOW STATUS LIKE 'Slow_queries';"
```

Table analysis

```
mysql -u slurm -p slurm_acct_db -e "SHOW TABLE STATUS;"
```

Key Takeaways

Recovery Principles

1. Systematic Diagnosis - logs → config → network → hardware
2. Service Dependencies - respect startup order
3. Backup Strategy - automate daily, test monthly
4. Change Control - test first, rollback ready
5. Monitoring - proactive alerts prevent outages

Most Critical Skills

- Log analysis - 80% of issues are in logs
- Config management - keep systems synchronized
- Database maintenance - accounting is business critical
- Network troubleshooting - foundation of cluster communication

Emergency Contacts

- Keep vendor support numbers handy
- Maintain escalation procedures
- Document tribal knowledge

Q&A & Discussion

Discussion Topics

- Site-specific experiences with Slurm failures
- Local backup/recovery procedures
- Integration with monitoring systems
- Automation opportunities

Resources

- Slurm Documentation: <https://slurm.schedmd.com/documentation.html>
- Troubleshooting Guide:
<https://slurm.schedmd.com/troubleshoot.html>
- Mailing Lists: slurm-users@lists.schedmd.com

Contact Information

- Internal escalation procedures
- Vendor support contacts
- Emergency response team

Bonus

Emergency Command Cheat Sheet

Service status

```
systemctl status slurmctld slurmdbd
```

```
sinfo -NeI # Node status
```

```
squeue -u $USER -o "%i %t %r" # Job status
```

Quick recovery

```
systemctl restart slurmctld
```

```
scontrol update NodeName=X State=resume Reason="fixed"
```

```
scontrol reconfig # Reload config
```

Emergency contacts

```
tail -100 /var/log/slurm/slurmctld.log
```

```
journalctl -u slurmctld -n 50
```

```
mysql -u slurm -p slurm_acct_db
```

Backup restore

```
systemctl stop slurmctld
```

```
cp /backup/slurmctld.state.YYYYMMDD /var/lib/slurm/slurmctld.state
```

```
systemctl start slurmctld
```

MPI/Parallel Execution Failures

MPI Version Mismatch - #1 User Issue

Error: "OPAL ERROR: Unreachable in file pmix3x_client.c"

Diagnosis

```
module list
```

```
which mpirun
```

```
ldd $(which mpirun) | grep pmix
```

```
srun --mpi=list
```

Resolution

```
module purge
```

```
module load gcc/12.1.0 openmpi/4.1.5
```

```
srun --mpi=pmix_v4 ./mpi_app
```

InfiniBand/Network Issues

Test connectivity

```
srun -N 2 hostname
```

```
srun -N 2 --ntasks=2 ib_write_bw
```

Force IB transport

```
export OMPI_MCA_btl=self,vader,openib
```

```
export UCX_NET_DEVICES=mlx5_0:1
```